

Aniket Kshirsagar

Data Analyst

New Jersey, USA | aniket.kshirsagar.work08@gmail.com | [LinkedIn](#) | [GitHub](#)

PROFESSIONAL SUMMARY

Data Analyst with 5+ years helping teams in banking, healthcare, and retail make better decisions from messy data. I start with the real business problem, choose the analysis that fits, and judge my work by whether it changed what someone did, with measurable outcomes: sharper fraud detection, faster reporting, and forecasts accurate enough to reshape planning.

SKILLS

- Programming:** Python, SQL
- Data Analysis and Manipulation:** Pandas, NumPy, Excel (Advanced), Power Query
- Visualization and BI:** Power BI, Tableau, Looker, Matplotlib, Seaborn, KPI Dashboards, Data Storytelling
- Databases:** MySQL, PostgreSQL, SQL Server, Oracle
- Statistics:** Hypothesis Testing, Regression Analysis, A/B Testing, Forecasting
- Cloud and Big Data:** AWS (S3, Redshift, Athena, Glue), Azure, Google BigQuery, Apache Spark (PySpark), Hadoop (HDFS, Hive)
- ETL:** Data Wrangling, ETL Pipelines, Data Transformation, Azure Data Factory, dbt
- Fraud Detection:** Ravelin, Anomaly Detection, Rule-Based Models
- Generative AI:** LLM-Based AI Agents (Build and Deploy), NLP
- Collaboration:** Git, GitHub, Jira, Confluence, Agile

WORK EXPERIENCE

Senior Data Analyst, JP Morgan Chase Mar 2026 to Present
New Jersey, USA

- Investigators were losing hours chasing false fraud alerts, so we traced exactly where the detection rules misfired and rebuilt them; accuracy climbed 28% and the team could finally spend its time on the cases that genuinely mattered.
- Built automated, standardized monitoring to replace the slow manual checks running across every banking channel, which cut investigation effort 22% and gave risk and compliance one reliable daily view they could act on with confidence.
- Leaders could not see fraud trends until the damage was already done, so we surfaced the anomalies early through clear, shared dashboards and helped the business cut fraud-related exposure by roughly 15% over the course of the year.

Full Stack Developer, Astoria AI (Part-Time) Sept 2025 to Feb 2026
New York, USA (Remote)

- Promising AI-agent prototypes were sitting unused because no one could run them, so we took them to production as reliable services that now automate real recruiting workflows from first contact through to the very final stage.
- Recruiters were stuck doing manual follow-up, so we built a pipeline that advances candidates automatically with scheduled reminders, removing most of the repetitive coordination they used to handle by hand every single day.
- Implemented a single dashboard to monitor and operate every agent in one place, all secured behind OAuth-protected APIs, which turned a loose set of scripts into a product the team could actually run and trust day to day.

Data Analyst Intern, JPMorgan Chase (Part-Time) May 2024 to Aug 2024
New York, USA (Remote)

- Analyzed slow-running SQL queries across production reporting workflows, using execution plans to pinpoint full table scans, inefficient joins, and missing predicates as the primary bottlenecks.
- Rewrote query logic by converting subqueries to joins, pushing filters earlier in execution, and removing redundant operations, cutting average query runtime materially for downstream reporting.
- Designed and implemented indexes (including composite and covering indexes) on high-traffic tables, reducing query cost and improving dashboard load times for business users.
- Built Tableau dashboards to stakeholder requirements, enabling self-service access to insights that previously required manual SQL pulls.

Data Analyst, Capgemini Jan 2022 to July 2023
India

- Healthcare reporting was manual and error-prone, so we automated the whole process end to end, cutting effort by nearly 28% and giving stakeholders numbers they could trust enough to act on without double-checking them first.
- Analysis kept breaking on fragmented source systems, so we consolidated them with Azure Data Factory into clean, analytics-ready data that the wider enterprise could genuinely depend on for its everyday reporting and decisions.
- Care teams needed to know who to prioritize, so we segmented patients and modeled cohorts in Python, pointing limited resources and interventions toward the highest-risk groups first instead of spreading them thin across everyone.

Junior Data Analyst, Accenture
India

Aug 2019 to Dec 2021

- Stakeholders needed to see what the data was actually telling them, so we built Tableau charts and dashboards to each requirement, turning raw datasets into clear, readable visuals the whole team could read and act on.
- Messy inputs were dragging down everything downstream, so we ran careful data validation, handling nulls and duplicates and trimming stray whitespace so the data feeding the models started clean, consistent, and trustworthy.
- That validated data fed the Python models our data engineers were building, and the cleaner the inputs we handed them, the more efficiently and accurately those models could run on top of all the work we had already done.

PROJECTS

Soccer PepStats, End-to-End Football Analytics Platform

[Live](#) | [Code](#)

- Python, pandas, DuckDB, FastAPI, React/TypeScript
- I wanted to answer real tactical questions in football, so I turned about 1.3M raw StatsBomb events into xG, shot maps, and player-season views inside a 12-tab interactive dashboard that anyone on a staff could open and explore.
- Built it on a batch layer plus a resilient real-time layer sitting behind a tested FastAPI service, so the numbers stay fast and genuinely trustworthy even when someone is making a real call under pressure during a live match.

COURTSIDE, End-to-End NBA Analytics Platform

[Live](#) | [Code](#)

- Python, DuckDB, dbt, React/TypeScript
- I wanted to understand basketball the way a front office does, so I turned 13.5M play-by-play events across 77 full seasons into advanced metrics and league-wide player percentiles that would actually hold up to real front-office scrutiny.
- Modeled 16 messy, unrelated sources into clean, automatically tested data marts and shipped a polished 6-page dashboard, carrying a single question all the way from raw data through to a decision-ready product in production.

SelfPrep, AI Voice-Agent Interview Practice Platform

[Live](#) | [Code](#)

- Next.js, TypeScript, Firebase, Vapi AI, Google Gemini, Tailwind
- Built a job-interview practice platform with Vapi AI voice agents: candidates run a spoken, adaptive interview, and Google Gemini generates the questions and a structured feedback report scored across technical knowledge, communication, confidence, and problem-solving.

GradEase, Graduation E-Commerce Platform (Pace University Capstone)

[Code](#)

- MERN stack, AI chatbot
- For my Pace University capstone, I built a full-stack MERN store paired with a Voiceflow AI chatbot that gave students personalized product recommendations and real-time order tracking the whole way through checkout.

EDUCATION

Master of Science in Computer Science, Pace University, New York, USA
GPA 3.5/4

May 2025

Bachelor of Science in Computer Science, Savitribai Phule University, Pune, India
GPA 3.4/4

Aug 2021

CERTIFICATIONS

- INSPIRE (International Student Professional Readiness Education) Program Certification, Pace University (Nov 2024)

ACHIEVEMENTS

- Held robotics leadership roles across all three years of the Bachelor's program, advancing from Technical Coordinator of the robotics team, to Technical Head of robotics events, and finally to Head of the flagship robotics event.

- State-level under-18 soccer player who went on to represent the college team across inter-college tournaments and play for the Computer department in a hard-fought inter-department final against a strong and determined ENTC department.